

STAND paper console

One page for both authors: where the paper stands, what is frozen, what runs this week, and what separates it from an accept. Repo: [anishesg/reality-monitoring](https://github.com/anishesg/reality-monitoring) (commit d88028c).

6 days

to full-paper deadline, Sept 25 11:59 PM AoE. Abstract submitted.

0 / 9

pre-registered runs finished. Nothing has trained yet.

0 vCPU

Azure GPU quota. Ticket 2609190040000347 open; queue script armed.

\$10,000

sponsorship credit, \$0 spent, kill switch tested.

17

checkpoints measured, 183k v17 trials, 4 figures built.

1 · ABSTRACT CLAIMS VERSUS WHAT THE REPOSITORY CAN BACK

What is proven, what is promised

CLAIM IN THE SUBMITTED ABSTRACT	EVIDENCE IN REPO	STATUS	OWNER
Naming an alternative flips as often as citing a source (0.92 vs 0.92, Llama-3.1-8B)	v17 cells, <code>analyze_v17.py</code>	in repo	Anish
Confidence AUROC 0.49→0.70 with scale; behavioral use ~0.05 vs truth effect 0.68–0.81	v16 + v17	in repo	Anish
Matched-position interaction −0.38, 95% CI [−0.51, −0.26]	<code>analyze_glmm.py</code> , family-level meta	in repo	Anish
Preference optimization raises capitulation in 2/3	v17 stage ladder (OLMo-2, Tulu-3, Zephyr)	in repo	Anish

lineages; RL does not reverse

Reliability note 90% vs 30% (0.91 vs 0.78, Qwen2.5-7B)	none	not pushed	Anish
SFT on demonstrations restores confidence use, at an accuracy cost	none	not pushed	Anish
Resistance to repeated challenges predicts correctness in 16/19	none (no multi-turn cell)	not pushed	Anish
Fresh-context re-ask recovers accuracy (1.00 vs 0.07)	none; also reads as tautological under greedy decoding	not pushed · reframe	Anish
"19 checkpoints, ~400,000 trials"	17 in v17, 19 in v16 (two failed); count only holds pooled across v15–v17	verify	both

Also on the submission: delete the literal `\end{abstract}` at the end, and rephrase "surface form rather than evidential content" to "the challenge's *stated* evidence", since the truth effect shows models do respond to content they can verify.

2 · FROZEN RUN LIST (PREREG, 2026-09-19)

What runs this week

RUN	WHAT	GPU-H	TESTS	STATUS
R1	OLMo-2-7B-SFT: A0 reuse, A2, A1, A3, A4, A5 on TEST	25–32	H1, H2, H3, H4, H5	queued · no GPU
R2	Tulu-3-8B-SFT: A0, A2, A1, A3 on TEST	~15	replication	queued

R3	A2 sensitivity sweep on DEV ($\beta \in \{.05, .1, .2\}$; $\text{lr} \in \{2\text{e-}6, 1\text{e-}5\}$); selection rule fixed	~8	hyperparameter provenance	queued
R4	LLM-judge agreement on 500 stratified TEST trials	0 (API)	grading validity	runs on laptop
R5	Stage measurement, published 13B and 32B SFT/DPO/Instruct	~12	H7 (scale)	queued
R6C	OLMo-2-32B-SFT: A0, A2 (4×80 GB)	~25	H8 (fix at scale)	conditional
R7C	OLMo-2-13B-SFT: A0, A2, A1, A3	~25	scale trend	conditional
R8C	32B A3, only if R1 A3 retain-correct ≥ 0.6	~120	H8	conditional
R9C	Frontier-API decomposition, 150 questions × self-origin cells, 4–5 API models	0 (API, ~\$200)	frontier relevance	runner to write

Splits: TEST = hard-bank 0–449 (all paper numbers). DEV = 800–899 (sweep only, never reported). TRAIN = SciQ + 450–799. Every hyperparameter is now tagged standard / heuristic / harness-fixed in the prereg; the 0.70 frontier line is a marker, the claim is Pareto dominance with CIs.

3 · WHAT RECENTLY ACCEPTED PAPERS HAVE THAT WE DO NOT YET

The acceptance gap

Read across sycophancy and alignment papers accepted at ICLR, NeurIPS, ICML and *ACL in the last two years (Persuasion-Balanced Training, pinpoint tuning, synthetic-data debiasing, and the 2026 pushback-robustness benchmarks). The pattern of what they carry, and where this paper stands.

FEATURE OF
ACCEPTED PAPERS

THIS PAPER TODAY

CLOSES WITH

Evaluates current frontier models, not only $\leq 14B$ open weights	missing	R9: same harness over APIs. Biggest single lever; the 0.92-vs-0.92 finding at the frontier is the quotable sentence.
Causal intervention, not correlation across released checkpoints	built, not run	R1 A1: train the DPO stage ourselves and show the pressure jump appear.
A fix that does not cost capability	built, not run	R1/R2 A3 with MMLU, GSM8K, IFEval within 1 point. Contrast with the SFT-at-an-accuracy-cost result already in the abstract.
Scale trend within the paper's own lineage	missing (14B ceiling)	R5 + R6: 7B $\rightarrow$ 13B $\rightarrow$ 32B on the same OLMo-2 recipe.
Grading validated against an LLM or human judge	missing	R4. Substring grading is known to undercount capitulation by 18–24 points; report agreement.
Pre-registration and honest negative results	present	Keep the failed self-vs-other prediction and the steering null in the paper; reviewers credit this.
Hyperparameter sensitivity reported	planned	R3 on DEV, appendix table.
Clear positioning against the nearest prior work	partial	Related-work paragraph naming PBT, Pressure-Tune, BTS-GRPO, Certainty Robustness, and the July 2026 steering null, with the three things none of them have: pressure-only challenges, source/content decomposition, stage-ladder causal replication.
Released artifacts match every claim	four claims unbacked	Anish pushes the four unpushed experiments with trial files, or they leave the abstract.

Vikram's deliverables

- A** R1 A1 on both backbones: the causal replication. Turns the paper's lineage finding into a controlled result.
- B** R1/R2 A2 and A3: the fix, with the capability table. If retention rises without accuracy loss, the abstract's last sentence becomes this result.
- C** R4 judge check and R3 sweep: the two methodological exposures a reviewer would raise first.
- D** R5/R6 scale rung and R9 frontier APIs: the two relevance gaps.
- E** Write the training methods section, the causal results section, the limitations, and own all figures (the four existing ones plus frontier-with-arrows and capability retention).

- 01 No GPU by Monday.** Azure quota is zero and auto-requests are rejected. Mitigation: della via Anish today; second Azure subscription (pay-as-you-go under the same billing profile draws the same credit); AWS On-Demand P quota request.
- 02 GRPO arms miss the deadline.** 6–9 h each, five per backbone. Ship with A1/A2 in the main text and GRPO as a partial table if not done by Tuesday night.
- 03 Fix costs capability or over-holds.** Report the frontier honestly; a 40-point retention gain with a small acceptance drop is still a strong paper.
- 04 Unbacked abstract claims.** Anything not in the repo by Thursday comes out of the abstract.
- 05 Scope creep.** The nine runs are frozen. Anything new goes to a follow-up.

6 • DAY BY DAY TO SEPT 25

The week

- | | |
|--------|--|
| Sat 19 | Anish pushes unpushed experiments. Calls: Microsoft (second subscription), AWS (P-instance quota). Anish starts R1 on della via the handoff prompt. Write the R9 API runner. |
| <hr/> | |
| Sun 20 | R3 sweep and R4 judge check. R9 frontier runs from a laptop. Draft Sections 1–3 and related work from existing results. |
| <hr/> | |

Mon 21

A2/A1 numbers on both backbones. Rerun A2 with the sweep-selected setting if it differs. Write Section 4 around A1 and A2. R5 stage measurement.

Tue 22

GRPO arms land. Frontier-with-arrows and capability figures. Decide R6/R7/R8. Rewrite the abstract to what landed.

Wed 23

Full draft with appendices: prereg verbatim, prompts, exclusion rates, judge agreement, per-checkpoint tables, negative results.

Thu 24

Internal review against the reviewer-threat list. Every abstract number checked against a table. Figures final.

Fri 25

Buffer. Submit before 11:59 PM AoE.

7 • FIGURES ALREADY IN THE PAPER

Measurement results

Mere mention drives most capitulation




Fig 1. Mere mention of an alternative drives most capitulation; source adds little.

The DPO stage installs pressure capitulation; RLVR does not repair it




Fig 3. The DPO stage installs pressure capitulation; RLVR does not repair it.




Fig 4. No checkpoint reaches the retain/accept frontier. The training arms will be drawn as arrows from A0.




Fig 2. Confidence becomes informative with scale; controlled use stays near zero.

Sources: repo state at d88028c; Azure subscription eae276d8; prereg files under prereg/ ; comparison papers: Stengel-Eskin et al. NAACL 2025, Chen et al. ICML 2024, Wei et al. 2023, Mytsyk et al. 2026, Saadat & Nemzer 2026, Aamir & Adil 2026.